



Vcube Exam — Answer Sheet

Vcube Software Solutions · Multi Cloud DevSecOps Program

STUDENT

Mohammed Abdul Sami

BATCH

AWS Batch 18

SCORE

23/25

RESULT

92% ·  Outstanding

#	QUESTION · YOUR ANSWER · CORRECT ANSWER · EXPLANATION	STATUS
1	Which of the following IP addresses belongs to a private (RFC 1918) range? Your Answer: 172.20.5.10 Correct Answer: 172.20.5.10  Explanation: 172.16.0.0/12 (172.16-172.31) is private. 8.8.8.8 and 1.1.1.1 are public DNS; 100.64.0.0/10 is carrier-grade NAT space, not RFC 1918.	✓ Correct
2	Which target types can you register behind a Network Load Balancer target group? Your Answer: EC2 instances, IP addresses, and even an Application Load Balancer Correct Answer: EC2 instances, IP addresses, and even an Application Load Balancer  Explanation: NLB supports instance and IP targets, and you can front an ALB with an NLB to combine static IPs with Layer 7 features.	✓ Correct
3	Approximately how many total IP addresses does a 10.0.0.0/16 network contain? Your Answer: 65,536 Correct Answer: 65,536  Explanation: A /16 leaves 16 bits for hosts: $2^{16} = 65,536$ addresses (the /16 VPC CIDR you see in AWS examples).	✓ Correct
4	In a hub-and-spoke AWS network design, which service acts as the central hub connecting many VPCs and on-premises networks? Your Answer: AWS Transit Gateway Correct Answer: AWS Transit Gateway  Explanation: Transit Gateway is the cloud router: each VPC and VPN attaches to the hub instead of building many peer-to-peer links.	✓ Correct
5	Before a VPC peering connection becomes active, what must happen? Your Answer: The owner of the acceptor VPC must accept the peering request Correct Answer: The owner of the acceptor VPC must accept the peering request  Explanation: Peering is a requester/accepter handshake; after acceptance you still must update route tables in both VPCs.	✓ Correct
6	By default, a newly created custom Security Group allows all outbound traffic. Your Answer: True Correct Answer: True  Explanation: New custom SGs start with no inbound rules and an outbound rule allowing 0.0.0.0/0; you then add inbound allows as needed.	✓ Correct

#	QUESTION · YOUR ANSWER · CORRECT ANSWER · EXPLANATION	STATUS
7	For which workload is AWS Global Accelerator often a better fit than CloudFront? Your Answer: Non-HTTP TCP/UDP applications such as gaming or IoT that need static entry IPs Correct Answer: Non-HTTP TCP/UDP applications such as gaming or IoT that need static entry IPs Explanation: Global Accelerator works at the network layer for any TCP/UDP traffic and gives static Anycast IPs; CloudFront is an HTTP/S caching CDN.	✓ Correct
8	AWS Config rules can automatically flag a resource that is non-compliant, for example an unencrypted S3 bucket. Your Answer: True Correct Answer: True Explanation: Rules mark resources COMPLIANT or NON_COMPLIANT and can trigger automatic remediation actions.	✓ Correct
9	Which ALB listener rule condition sends api.example.com traffic to one target group and www.example.com to another? Your Answer: Host-header (host-based) routing Correct Answer: Host-header (host-based) routing Explanation: Host-header rules look at the requested domain; path-pattern rules look at the URL path (e.g. /api/*).	✓ Correct
10	Why is CloudTrail valuable for security and auditing of an AWS account? Your Answer: The CloudTrail is valuable for security and auditing of an AWS account, it stores all the metrics logs actions performed by an individual user no matter in the creating the VPC, IAM policies, IAM roles, and EC2 network configuration and etc. This helps security teams in identifying the criminal or any security incident is occurred like data breaches and etc. Correct Answer: CloudTrail records every API call with the caller's identity, time and source, creating an audit trail. Security teams use it to investigate incidents and unauthorized changes, and auditors use it as evidence for compliance. Explanation: Without CloudTrail you cannot answer 'who deleted/changed this resource and when'.	✗ Wrong
11	What is an IAM policy? Your Answer: A JSON document that defines permissions - which actions are allowed or denied on which resources Correct Answer: A JSON document that defines permissions - which actions are allowed or denied on which resources Explanation: Policies are attached to users, groups and roles; they state Effect (Allow/Deny), Action and Resource.	✓ Correct
12	What does an EC2 Auto Scaling Group do? Your Answer: Automatically adds or removes EC2 instances to maintain desired capacity and replace unhealthy instances Correct Answer: Automatically adds or removes EC2 instances to maintain desired capacity and replace unhealthy instances Explanation: ASGs keep instance count between min/max, follow scaling policies, and swap out instances that fail health checks.	✓ Correct

#	QUESTION · YOUR ANSWER · CORRECT ANSWER · EXPLANATION	STATUS
13	What does AWS Config do? Your Answer: Continuously records the configuration of your AWS resources and evaluates them against desired rules for compliance Correct Answer: Continuously records the configuration of your AWS resources and evaluates them against desired rules for compliance  Explanation: Config keeps a configuration history per resource and runs managed/custom rules (e.g. 'is encryption enabled?') against them.	✓ Correct
14	What does an AWS Site-to-Site VPN connection do? Your Answer: Creates an encrypted IPsec tunnel over the internet between your on-premises network and your VPC Correct Answer: Creates an encrypted IPsec tunnel over the internet between your on-premises network and your VPC  Explanation: Site-to-Site VPN links your data center (customer gateway) to AWS (virtual private gateway or Transit Gateway) with IPsec encryption.	✓ Correct
15	In an Auto Scaling Group, what does a launch template specify? Your Answer: The configuration for new instances: AMI, instance type, key pair, security groups and user data Correct Answer: The configuration for new instances: AMI, instance type, key pair, security groups and user data  Explanation: Launch templates (or the older launch configurations) tell the ASG exactly how to build each new instance.	✓ Correct
16	What is CloudWatch Logs used for? Your Answer: Centralizing and storing log data from EC2, Lambda and other AWS services for search and analysis Correct Answer: Centralizing and storing log data from EC2, Lambda and other AWS services for search and analysis  Explanation: Logs from many services stream into log groups; you can filter, query with Logs Insights, and set metric filters on patterns.	✓ Correct
17	An Auto Scaling Group can launch instances across multiple Availability Zones for better availability. Your Answer: True Correct Answer: True  Explanation: You attach the ASG to subnets in several AZs and it balances instances across them.	✓ Correct
18	What does AWS CloudTrail primarily record? Your Answer: API calls and account activity - who did what, when, and from where Correct Answer: API calls and account activity - who did what, when, and from where  Explanation: CloudTrail logs management (and optionally data) events with identity, timestamp and source IP - the audit history of the account.	✓ Correct
19	Which S3 feature automatically moves objects to cheaper storage classes (for example to Glacier) based on rules you define? Your Answer: S3 Lifecycle policies Correct Answer: S3 Lifecycle policies  Explanation: Lifecycle rules transition objects (Standard -> IA -> Glacier) or expire them after the age you specify.	✓ Correct

#	QUESTION · YOUR ANSWER · CORRECT ANSWER · EXPLANATION	STATUS
20	What is a Multi-AZ deployment in Amazon RDS and why would you enable it? Your Answer: A Multi-AZ deployment in Amazon RDS is the secondary Database server in the other or multi Availability Zones, which acts as a backup when we the primary database is down or failover due to some availability issue. In the Multi-AZ deployment, the Database servers are created in the different Availability Zone in case of Data center connectivity issue due to wars etc. As much as I remember while doing the hands-on. There are two type of Multi-AZ deployment in Amazon RDS, 1. 3 level Multi-AZ deployment - it creates two secondary databases (backups) under Multi-AZ deployment in 2 different availability zones and if primary database is down, it stores in the two secondary databases as per their activeness 2. 2 level Multi-AZ deployment - it create one secondary database (backup) under Multi-AZ deployment in 1 different availability zones and if primary database is down it stores in the secondary databases. The secondary databases server is usually in the readonly mode when primary server is active and available (Read/Write mode). Once the primary serve isn't available, the secondary database mode becomes read/write mode. Correct Answer: Multi-AZ keeps a synchronous standby replica of the database in a different Availability Zone. If the primary or its AZ fails, RDS automatically fails over to the standby with minimal downtime, improving availability and durability (it is not for read scaling - that is what read replicas are for). Explanation: Multi-AZ = high availability/automatic failover; read replicas = scaling reads.	✗ Wrong
21	A CloudWatch alarm can trigger EC2 Auto Scaling actions, for example adding instances when average CPU is high. Your Answer: True Correct Answer: True Explanation: Target/step scaling policies are driven by CloudWatch alarms on metrics like CPUUtilization.	✓ Correct
22	An interface VPC endpoint connects your VPC privately to a supported service using which technology? Your Answer: AWS PrivateLink with an elastic network interface inside your subnet Correct Answer: AWS PrivateLink with an elastic network interface inside your subnet Explanation: Interface endpoints are ENIs with private IPs powered by PrivateLink; gateway endpoints (S3, DynamoDB) are the other endpoint type.	✓ Correct
23	Which CloudWatch feature watches a metric and notifies you or triggers an action when it crosses a threshold? Your Answer: CloudWatch Alarms Correct Answer: CloudWatch Alarms Explanation: Alarms fire on thresholds (e.g. CPUUtilization > 80%) and can send SNS notifications or trigger Auto Scaling.	✓ Correct
24	Which EBS volume type is a general-purpose SSD offering a balance of price and performance (and is commonly the default)? Your Answer: gp3 (General Purpose SSD) Correct Answer: gp3 (General Purpose SSD) Explanation: gp3 suits most workloads with baseline 3,000 IOPS and 125 MB/s; io2 targets high-performance databases; st1/sc1 are HDD.	✓ Correct

#	QUESTION · YOUR ANSWER · CORRECT ANSWER · EXPLANATION	STATUS
25	In a Site-to-Site VPN, what is the role of the Virtual Private Gateway (VGW)? Your Answer: The role of the Virtual Private Gateway is that it creates a secure tunnel over the Internet, where two or more devices can communicate with each other. But, in our case, we have created VGW to connect the machines from the Azure VNG. While connecting the VGW, we configured the Site to Site VPN where we have provided the Azure Vnet IP address CIDR range. Where we actually create the VPN for the connection to Private Internal IP address of Azure VM to this AWS EC2 Machine And In the end, we add this in the Route Table and select Virtual Private Gateway for the connection. And we were able to ping or communicate between AWS EC2 instance and Azure VM. Correct Answer: The Virtual Private Gateway is the VPN concentrator on the AWS side. It is attached to your VPC, terminates the two IPsec tunnels coming from your on-premises customer gateway, and routes the encrypted traffic into the VPC.  Explanation: The pair is: customer gateway on-premise, VGW on AWS, with two tunnels for redundancy between them.	 Correct

Generated by Vcube Exam Portal · 19/09/2026, 12:25:13